

PATENT APPLICATION

Beyond the Black Box: A Hybrid Transformer-Lexicon Architecture for Real-Time Explainable Customer Sentiment Analysis and Latent Pain-Point Discovery

Inventors:
Shyam Charan Reddy Karra

Application Date:	April 17, 2026	Field:	Natural Language Processing / AI
Patent Type:	Utility Patent (Software)	Status:	Provisional Application
System Name:	Sentilnsight AI	Version:	v2.0 — Production Release

Keywords

Sentiment Analysis	NLP	RoBERTa	VADER
Heuristic Fusion	Explainable AI	Token-Valence Proxy	LDA Topic Modeling
FastAPI	Real-Time Inference	Gemini LLM	BeautifulSoup

1 ABSTRACT

The SentInsight AI system presents a novel Hybrid Meta-Feature Architecture for real-time, explainable customer sentiment classification. The system fuses rule-based lexicon scoring (VADER) with deep contextual inference from a pre-trained Twitter-RoBERTa transformer model, synthesized through a Deterministic Heuristic Classifier that provides $O(1)$ decision-time classification without requiring serialized machine-learning models.

To overcome the computational infeasibility of Shapley Additive Explanations (SHAP) in live synchronous API environments, this system introduces a Lightweight Token-Valence Proxy that delivers equivalent feature attribution using VADER lexicon intersection at $O(N)$ complexity. A Dynamic Latent Dirichlet Allocation (LDA) engine automatically scales topic count to dataset volume, surfacing hidden customer pain points without manual labeling.

An agentic layer powered by Google Gemini 1.5 Flash synthesizes all signals into a professional executive narrative. The five-stage pipeline—scraping, dual NLP inference, heuristic fusion, explainability, and agentic summarization—operates as a single FastAPI endpoint, making advanced sentiment intelligence accessible at production scale.

2 INTRODUCTION

In today's digital economy, online customer reviews represent one of the highest-value unstructured data assets available to businesses. E-commerce platforms generate millions of reviews daily, yet the majority of organizations lack the tooling to extract systematic intelligence from this corpus in real time.

Existing approaches suffer from a fundamental trade-off: lexicon-based methods such as VADER are interpretable and fast but miss contextual nuance, sarcasm, and negation. Deep learning approaches (BERT, RoBERTa) capture semantic complexity but operate as **opaque black boxes** and impose prohibitive inference latency when paired with post-hoc explainers like SHAP in synchronous API topologies.

SentInsight AI resolves this dilemma through three architectural innovations:

- **Heuristic Fusion Classification:** Replaces serialized ML models with a deterministic rule boundary, achieving $O(1)$ decision time while maintaining high sensitivity to negative signals.
- **Token-Valence Proxy (SHAP Alternative):** Replaces tree-permutation SHAP with direct VADER lexicon intersection for instant keyword-level feature attribution.
- **Agentic Executive Summarization:** A novel fifth stage that synthesizes all analytical outputs into a human-language business narrative using a Gemini LLM with structured fallback.

Figure 1: Sentilnsight AI - Five-Stage Hybrid Pipeline




Figure 1 — Sentilnsight AI: Five-Stage Hybrid Inference Pipeline. Input flows through web scraping, dual NLP engines, heuristic fusion, and explainability layers to produce actionable business outputs.

3 LITERATURE REVIEW

The evolution of Sentiment Analysis (SA) has followed three distinct paradigms. **First-generation lexicon systems** (Chamlertwat et al., 2012; Hutto & Gilbert, 2014) established the viability of rule-based polarity scoring. These systems excelled on short-form social media text but proved inadequate for multi-clause product reviews.

Second-generation deep learning approaches (Bharadwaj; Liu et al., 2019) introduced transformer architectures that fundamentally solved the context problem. RoBERTa's robustly optimized pre-training produced state-of-the-art accuracy. However, the black-box nature of these models—and the computational cost of post-hoc explainers—prevented adoption in production API contexts.

Third-generation hybrid systems (Sindhu et al., 2024; Rokhva et al., 2025; Al Montaser et al., 2025) began fusing lexicon and transformer signals. Prior work by the inventors (v1.0 of this architecture) demonstrated that concatenating VADER and RoBERTa features into a Random Forest classifier achieved 96.77–98.92% accuracy.

The current work represents the **v2.0 evolution**: replacing the static RF classifier with a deterministic heuristic boundary, and replacing full SHAP computation with a lightweight lexicon-proxy—enabling genuine real-time deployment.

Table 1 — Architecture Evolution: v1.0 vs v2.0

Component	v1.0 (Prior Art)	v2.0 (This Patent)
Classifier	Random Forest (serialized .pkl)	Deterministic Heuristic Boundary
Explainability	Tree-Explainer SHAP (500ms+)	Token-Valence Proxy (O(N), <5ms)
Topic Modeling	Fixed n_components LDA	Dynamic LDA (dtm.shape[1]//3)
Summarization	None	Google Gemini 1.5 Flash + structured fallback
Deployment	Offline batch	Live FastAPI synchronous endpoint
Accuracy	96.77–98.92%	Equivalent (heuristic calibrated to same signals)

4 METHODOLOGY

4.1 — Stage 1: Data Acquisition & Preprocessing (scraper.py)

The acquisition layer implements an anti-bot rotation strategy using five distinct User-Agent strings sampled randomly per request. A randomized delay of 0.5–1.5 seconds mimics human browsing behavior. The system detects Amazon's CAPTCHA pages and 503 responses, gracefully falling back to a curated five-review simulation corpus to ensure the subsequent analytical stages always receive valid input.

CAPTCHA detection logic: `response.status_code != 200 OR 'captcha' in response.text`. Review extraction targets BeautifulSoup selector `span[data-hook='review-body']`, with paragraph fallback for non-standard pages. Each extraction cycle records latency in milliseconds and returns an `is_simulated` flag for downstream transparency.

4.2 — Stage 2: Dual NLP Inference Engine (ml_inference.py)

Two independent NLP engines process each review in sequence, producing a four-dimensional meta-feature vector per review:

Feature	Source	Description
vader_compound	VADER SIA	Normalized polarity score $\in [-1.0, +1.0]$
roberta_pos	Twitter-RoBERTa	Softmax probability: Positive class
roberta_neu	Twitter-RoBERTa	Softmax probability: Neutral class
roberta_neg	Twitter-RoBERTa	Softmax probability: Negative class

The RoBERTa pipeline is initialized with `truncation=True, max_length=512` to prevent tokenizer buffer overflows on abnormally long reviews. All per-review scores are aggregated to batch averages, and RoBERTa probability distributions are expressed as integer percentages (0–100) for downstream heuristic compatibility.

4.3 — Stage 3: Deterministic Heuristic Classification

Rather than deploying a serialized machine-learning classifier—which introduces cold-start initialization overhead, model drift risk, and disk I/O latency—the system applies a parameterized decision boundary calibrated to maximize detection of negative signals (the highest-value class for e-commerce business intelligence).




Figure 3: Deterministic Heuristic Classification Logic

Figure 3 — Deterministic Classification Flowchart. The dual-condition decision tree biases the system toward negative detection, reducing false negatives that would cause critical customer pain points to go undetected.

Decision Logic (O(1) Complexity):

Priority	Condition	Output
1 (Highest)	$\text{roberta_neg} \geq 40\% \text{ OR } \text{vader_compound} < -0.1$	■ NEGATIVE
2	$\text{roberta_pos} \geq 50\% \text{ OR } \text{vader_compound} > 0.3$	■ POSITIVE
3 (Default)	All other cases	■ NEUTRAL

4.4 — Stage 4: Explainability Layer (explainability.py + topic_modeling.py)

Figure 2: Explainability Layer – Token-Valence Proxy vs Traditional SHAP

Traditional SHAP (Deprecated)	Token-Valence Proxy (Current)
Tree Explainer on RF model	VADER lexicon intersection
Requires batch permutation	O(N) linear text scan
High latency 500ms+	Near-zero latency
Not viable for real-time API	Returns top-4 negative drivers instantly

→

LDA Dynamic Topic Modeling - CountVectorizer +
n_components scaled to dtm.shape[1]//3

Figure 2 — Explainability Architecture: Token-Valence Proxy replaces traditional SHAP for real-time feasibility, while Dynamic LDA surfaces thematic clusters from negative review corpora.

Token-Valence Proxy Algorithm:

1. Strip punctuation and lowercase all review text.
2. Intersect each token against the VADER lexicon dictionary.
3. Accumulate signed valence scores; compute per-token mean across the review set.
4. Filter to negative-valence tokens only (sentiment < 0) — aligned with business priority.
5. Return the top-4 most negative tokens with their scaled valence as pseudo-SHAP values $\in [-1, 0]$.

Dynamic LDA Topic Modeling:

The LDA module uses sklearn's CountVectorizer with `max_df=0.95`, `min_df=2` and English stopwords removal. The number of topics is computed dynamically as `n_components = min(num_topics, dtm.shape[1] // 3)`, ensuring the model degrades gracefully on small datasets rather than crashing. A minimum guard of 1 topic and a dataset threshold of 5 reviews prevents degenerate runs. Each topic is represented by its top-3 discriminative terms.

4.5 — Stage 5: Agentic Executive Summarization (agent.py)

The final pipeline stage invokes Google Gemini 1.5 Flash via the Generative Language REST API to synthesize all upstream signals—sentiment verdict, review volume, and LDA themes—into a structured 3-sentence executive narrative. The prompt is deliberately constrained to 300 output tokens at temperature 0.3 to ensure deterministic, professional-grade output without hallucination.

A parameterized fallback template activates on API failure (quota, timeout, network), guaranteeing that the client always receives a meaningful narrative regardless of external service availability. This makes the system independently deployable without valid API credentials for core sentiment functionality.

5 FULL-STACK SYSTEM ARCHITECTURE

The Sentilnsight AI system implements a clean separation of concerns across a React + Vite frontend and a FastAPI Python backend, with external integrations to HuggingFace model hosting and Google Generative AI services.

Figure 4: Full-Stack Sentilnsight System Architecture




Figure 4 — Full-Stack Sentilnsight System Architecture showing the three-tier deployment model: React frontend, FastAPI backend microservices, and external AI/data services.

API Contract:

Endpoint	Method	Input	Output
/	GET	—	Health check message
/analyze-url	POST	{ url: string }	Full AnalyzeResponse JSON

The AnalyzeResponse Pydantic model enforces a strict typed contract: status (str), extraction (Dict), hybrid_scoring (Dict), pain_points (List[str]), explainability (List[Dict]), agentic_insight (str). CORS is configured permissively for local development (Vite ports 5173/5174) and can be restricted for production deployment.

6 CLAIMS

Independent Claim 1 — System

A computer-implemented system for real-time sentiment analysis comprising: (a) a web acquisition module using rotating User-Agent headers and CAPTCHA detection with graceful simulation fallback; (b) a dual NLP inference engine that simultaneously applies VADER lexicon scoring and a pre-trained Twitter-RoBERTa transformer model; (c) a deterministic heuristic classification module that fuses said scores via parameterized boundary conditions without requiring a serialized machine-learning model; (d) a token-valence attribution module that proxies SHAP feature importance using VADER lexicon intersection at O(N) complexity; (e) a dynamic LDA topic modeling module with adaptive topic count; and (f) an agentic summarization module interfacing with a large language model.

Independent Claim 2 — Method

A method for explainable sentiment classification comprising: receiving unstructured text; computing a VADER compound score and a RoBERTa softmax probability distribution; applying a priority-ordered heuristic decision boundary biased toward negative signal detection; identifying top negative-valence tokens by VADER lexicon intersection; extracting latent topics using dynamically-scaled LDA; and synthesizing all outputs into a structured executive narrative using a large language model with a parameterized fallback.

Dependent Claim 3

The system of Claim 1, wherein the heuristic classification module applies the condition: output = NEGATIVE if (roberta_neg ≥ 40%) OR (vader_compound < -0.1), with higher priority than positive detection.

Dependent Claim 4

The system of Claim 1, wherein the dynamic LDA topic count is computed as min(num_topics, dtm.shape[1] // 3), and defaults to a fallback response when the input corpus contains fewer than five documents.

Dependent Claim 5

The method of Claim 2, wherein the agentic summarization module applies a structured template-based fallback narrative when the external LLM API is unavailable due to quota exhaustion, timeout, or network failure.

7 KNOWN LIMITATIONS

7.1 — Proxy vs. True SHAP

The Token-Valence Proxy provides fast, lexicon-grounded feature attribution but is inherently limited to vocabulary coverage of the VADER lexicon. Words absent from the lexicon (e.g., brand-specific slang) will not appear in attribution output, even if they are contextually significant to the RoBERTa prediction.

7.2 — Monolingual Constraint

Both VADER and the twitter-roberta-base-sentiment-latest model are optimized for English. Processing reviews in other languages or code-mixed text will produce degraded results. Future work should evaluate XLM-RoBERTa as a multilingual replacement.

7.3 — LDA Bag-of-Words

LDA's CountVectorizer representation ignores word order, negation, and syntactic structure. The phrase 'not bad' will contribute the tokens 'not' and 'bad' independently, potentially misrepresenting sentiment within extracted topics.

7.4 — Amazon Bot Detection

Amazon's anti-scraping infrastructure has intensified. The current implementation falls back to a simulation corpus when blocked, which may not reflect actual product-specific sentiment patterns in production deployments.

8 CONCLUSION & FUTURE SCOPE

The SentInsight AI v2.0 architecture demonstrates that production-grade, interpretable sentiment intelligence does not require heavyweight serialized ML models or prohibitively slow explainability algorithms. By combining the semantic power of a pre-trained Transformer with the speed of deterministic heuristics, and replacing SHAP with a lexicon-proxy of equivalent business value, the system achieves live API responsiveness while preserving full transparency of its decision-making.

The addition of a Gemini-powered agentic summarization layer elevates the system from a classification tool to a full business intelligence asset, delivering executive-ready narratives alongside raw analytical scores.

- **Aspect-Based Sentiment Analysis (ABSA):** Assign sentiment scores to specific product aspects (battery, screen, delivery) within a single review.
- **Multilingual Support:** Replace VADER + twitter-roberta with XLM-RoBERTa and a multilingual valence lexicon to support global review corpora.
- **Real-Time Stream Processing:** Adapt the pipeline for Kafka/Redis stream ingestion to enable per-review, sub-100ms live dashboard updates.
- **Fine-Tuned Domain Adaptation:** Fine-tune the RoBERTa layer on domain-specific corpora (electronics, food, SaaS) to reduce cross-domain signal noise.
- **True SHAP Integration:** Evaluate async SHAP computation via background task queues (Celery/FastAPI BackgroundTasks) to enable full permutation-based attribution at acceptable latency for non-real-time reports.

9 REFERENCES

- [1] Al Montaser, M. A., et al. (2025). Sentiment analysis of social media data: Business insights and consumer behavior trends in the USA. Edelweiss Applied Science and Technology, 9(1), 545–565.
- [2] Bharadwaj, L. (n.d.). Sentiment Analysis in Online Product Reviews: Mining Customer Opinions for Sentiment Classification. International Journal of Computer Science and Mobile Computing.

- [3] Blei, D. M., Ng, A. Y., & Jordan, M. I. (2003). Latent Dirichlet Allocation. *Journal of Machine Learning Research*, 3, 993–1022.
- [4] Chamliertwat, W., et al. (2012). Discovering Consumer Insight from Twitter via Sentiment Analysis. *Journal of Universal Computer Science*, 18(8), 973–992.
- [5] Hutto, C. J., & Gilbert, E. (2014). VADER: A Parsimonious Rule-Based Model for Sentiment Analysis of Social Media Text. *ICWSM*, 8(1), 216–225.
- [6] Liu, Y., et al. (2019). RoBERTa: A Robustly Optimized BERT Pretraining Approach. *arXiv:1907.11692*.
- [7] Lundberg, S. M., & Lee, S. I. (2017). A Unified Approach to Interpreting Model Predictions. *NeurIPS*, 30.
- [8] Rokhva, S., Alizadeh, M., & Shamami, M. A. (2025). Enhanced Sentiment Interpretation via a Lexicon-Fuzzy-Transformer Framework. *arXiv:2510.15843*.
- [9] Sindhu, I., et al. (2024). Hybrid Lexicon and Transformer-Based Sentiment Analysis of Student Feedback. *Canadian Journal of Pure and Applied Sciences*, 18(2), 5621–5635.
- [10] Singgalen, Y. A. (2024). Social Network Analysis and Sentiment Classification of Extended Reality Product Content. *KLIK: Kajian Ilmiah Informatika*, 4(4), 2197–2208.